

Numerical optimization for large scale problems and stochastic optimization

Constrained optimization

Elisa Salvadori (302630), Sonia Vittone(302673)

Contents

1	Quadratic Programming with Equality Constraints	2
2	Methods for the full problem	2
3	Methods for reduced form of the problem	3
3.1	Schur Complement Resolution	3
3.2	Null Space Method	3
4	Analysis of the problem	3
5	Testing and results	4
5.1	Case 1	4
5.2	Case 2	4
5.3	Case 3	5
5.4	Case 4	5
6	Conclusion	6
	Appendix	7

1 Quadratic Programming with Equality Constraints

The general problem of constrained optimization is written as:

$$\min_{x \in X} f(x),$$

where $f : \mathbb{R}^n \rightarrow \mathbb{R}$ continuously differentiable, and X is a nonempty, convex and closed set. We focus on the case of $f(x)$ quadratic with equality constraints only, we have:

$$\begin{aligned} \min f(x) \\ \text{s.t. } h(x) = 0, \end{aligned}$$

with $h(x)$ linear.

In particular, let us consider a Quadratic Programming problem with equality constraints:

$$\begin{aligned} \min \quad & \frac{1}{2}x^T Qx + c^T x \\ \text{s.t.} \quad & Ax = b, \end{aligned} \tag{1}$$

where $Q \in \mathbb{R}^{n \times n}$ is symmetric and positive (at least) semidefinite, $c \in \mathbb{R}^n$, $x \in \mathbb{R}^n$, $A \in \mathbb{R}^{m \times n}$ and $b \in \mathbb{R}^m$.

Now we want to apply the KKT conditions to (1), the Lagrangian function is:

$$L(x, \lambda) = \frac{1}{2}x^T Qx + c^T x + \lambda^T (Ax - b).$$

Recalling that if x^* is a local minimum for (1) and $\{\nabla h_i(x^*)\}_i$ are linearly independent vectors, then $\exists \lambda \in \mathbb{R}^m$ such that

$$\nabla f(x^*) + \sum_{i=1}^m \lambda_i \nabla h_i(x^*),$$

where λ_i are called Lagrange multipliers, so one necessary condition is that $\{\nabla h_i(x^*)\}_{i=1, \dots, m}$ are required to be linearly independent, this happens if rows of A are linearly independent that means A is full rank.

The KKT conditions (to be solved w.r.t. (x, λ)) for this kind of problem are the following:

$$\begin{cases} \nabla_x \mathcal{L}(x, \lambda) = Qx + c + A^T \lambda = 0 \\ \nabla_\lambda \mathcal{L}(x, \lambda) = Ax - b = 0 \end{cases} \tag{2}$$

Then we can use a matrix representation of the type:

$$\begin{bmatrix} Q & A^T \\ A & 0 \end{bmatrix} \begin{bmatrix} x \\ \lambda \end{bmatrix} = \begin{bmatrix} -c \\ b \end{bmatrix} \quad \implies \quad K\omega = d, \tag{3}$$

where K is again a symmetric matrix.

Now we can solve this new system considering the full problem or considering a reduce form of it.

2 Methods for the full problem

To solve the KKT conditions we could consider the full problem $K\omega = d$, with K symmetric and attack it:

- LDL^T factorization or LU factorization, but it could be problematic if $n + m$ is very large and K is sparse (fill-in problem);
- GMRES or other iterative methods.

3 Methods for reduced form of the problem

These methods consider reduced forms of $K\omega = d$:

- Schur Complement Resolution (Q non singular);
- Null Space Method.

3.1 Schur Complement Resolution

Starting from (2), from the first equation we obtain:

$$x = -Q^{-1}A^T\lambda - Q^{-1}c,$$

replacing it in the second equation, λ^* is the solution of:

$$\hat{Q}\lambda = -AQ^{-1}c - b,$$

where $\hat{Q} = (AQ^{-1}A^T)$ is the Schur complement (symmetric positive definite, assuming Q non singular and A full rank). Once we get λ^* we find x^* .

This approach is convenient if Q is easy to invert and if m is small w.r.t. n .

3.2 Null Space Method

Assuming that we know a particular solution \hat{x} to $Ax = b$ and we have a full rank matrix $Z \in \mathbb{R}^{n \times (n-m)}$ such that $AZ = 0_{m \times (n-m)}$ ($\dim(Ker(A)) = n - m$), a general solution to $Ax = b$ can be written as:

$$x = Z\nu + \hat{x}.$$

Replacing this solution into the first equation of (2):

$$Z^T Q Z \nu = -Z^T c - Z^T Q \hat{x},$$

solving the system for ν then we obtain x .

4 Analysis of the problem

Consider the problem

$$\begin{aligned} \min_{x \in \mathbb{R}^n} \quad & \sum_{i=1}^n x_i^2 - \sum_{i=1}^{n-1} x_i x_{i+1} + \sum_{i=1}^n x_i \\ \text{s.t.} \quad & \text{the sum } x_1 + x_{1+k} + x_{1+2k} + \dots = 1 \\ & \text{the sum } x_2 + x_{2+k} + x_{2+2k} + \dots = 1 \\ & \vdots \\ & \text{the sum } x_k + x_{2k} + x_{3k} + \dots = 1. \end{aligned}$$

This is a QP problem with:

- the matrix $Q \in \mathbb{R}^{n \times n}$ is tridiagonal with 2 in the main diagonal and -1 in the upper and in the lower diagonal,

$$Q = \begin{bmatrix} 2 & -1 & 0 & 0 & 0 & \cdots & 0 \\ -1 & 2 & -1 & 0 & 0 & \cdots & 0 \\ 0 & -1 & 2 & -1 & 0 & \cdots & 0 \\ \vdots & \ddots & \ddots & \ddots & \ddots & \ddots & \vdots \\ 0 & \cdots & 0 & -1 & 2 & -1 & 0 \\ 0 & \cdots & 0 & 0 & -1 & 2 & -1 \\ 0 & \cdots & 0 & 0 & 0 & -1 & 2 \end{bmatrix};$$

- the vector $c \in \mathbb{R}^n$ equal to $c = (1, 1, \dots, 1)^T$;
- the matrix $A \in \mathbb{R}^{k \times n}$ is formed by all identity matrix of dimension $k \times k$ repeated n/k times,

$$A = \begin{bmatrix} I_{k \times k} & I_{k \times k} & \cdots & I_{k \times k} \end{bmatrix} = \begin{bmatrix} 1 & 0 & \cdots & 0 & \cdots & 1 & 0 & \cdots & 0 \\ 0 & 1 & \cdots & 0 & \cdots & 0 & 1 & \cdots & 0 \\ \vdots & & \ddots & & \cdots & & & \ddots & \\ 0 & 0 & \cdots & 1 & \cdots & 0 & 0 & \cdots & 1 \end{bmatrix};$$

- and the vector $b \in \mathbb{R}^k$ as follows: $b = (1, 1, \dots, 1)^T$.

5 Testing and results

In this section we show our computational results for the problem using full form methods and reduced form methods. In particular we use LU and LDL' decomposition as direct methods and GMRES with tolerance equal to 10^{-12} and 100 iterations as the iterative one.

The problem is also tested with the Schur resolution, thanks to the non singularity of Q , with the inverse matrix (calculated directly using *inv* in MATLAB) and without it, and also with the Null Space method. To test the problem we use combinations of n and k , with $n = 10^4, 10^5$ and $k = 100, 500$.

5.1 Case 1

This case is tested with $n = 10^4$ and $k = 100$. The table shows the value of the function and the *Time* in seconds.

$n = 10^4, k = 100$											
LU		LDL'		GMRES		Schur		Schur autoinv		Null space	
$f(x)$	<i>Time</i>	$f(x)$	<i>Time</i>	$f(x)$	<i>Time</i>	$f(x)$	<i>Time</i>	$f(x)$	<i>Time</i>	$f(x)$	<i>Time</i>
100.00	0.5165	100.00	0.2023	100.00	0.3460	100.00	0.6746	99.99	0.0200	100	0.1737

Table 1: Comparison between the methods.

First we must consider that the iterative method (GMRES) doesn't converge to the desired tolerance $1e - 12$, the iterate returned has relative residual $1.2e - 06$.

The LDL' decomposition is faster than LU because it exploits the symmetry of K .

In general the reduced forms take less time than the full forms except for the Schur with inverse matrix, the latter takes longer time because it needs to invert the matrix Q that is not easily invertible. Schur autoinv which implements the *backslash* command is more efficient. Furthermore we can easily see that the solution value function found is almost the same for each method.

5.2 Case 2

This case is tested with $n = 10^4$ and $k = 500$. The table shows the value of the function and the *Time* in seconds.

$n = 10^4, k = 500$											
LU		LDL'		GMRES		Schur		Schur autoinv		Null space	
$f(x)$	<i>Time</i>	$f(x)$	<i>Time</i>	$f(x)$	<i>Time</i>	$f(x)$	<i>Time</i>	$f(x)$	<i>Time</i>	$f(x)$	<i>Time</i>
500.00	4.9963	500.00	0.4844	500.00	0.4937	500.00	1.1301	499.99	0.1073	500	0.0640

Table 2: Comparison between the methods.

The table underlines the same behaviour of the previous case, overall the full form methods take longer time than the reduced form methods. We must consider that the iterative method (GMRES) doesn't converge to the desired tolerance $1e-12$, the iterate returned has relative residual $6e-06$. The LDL' decomposition is faster than LU because it exploits the symmetry of K . Furthermore we can easily see that the solution value function found is almost the same for each method.

5.3 Case 3

This case is tested with $n = 10^5$ and $k = 100$. The table shows the value of the function and the *Time* in seconds.

$n = 10^5, k = 100$											
LU		LDL'		GMRES		Schur		Schur autoinv		Null space	
$f(x)$	<i>Time</i>	$f(x)$	<i>Time</i>	$f(x)$	<i>Time</i>	$f(x)$	<i>Time</i>	$f(x)$	<i>Time</i>	$f(x)$	<i>Time</i>
100.00	20.4642	100.00	7.6577	100.00	1.8866	Out of mem- ory	/	102.70	0.2223	100	100.8225

Table 3: Comparison between the methods.

For this case the matrix Q has dimension $10^5 \times 10^5$ and MATLAB can't compute the inverse matrix because it's too large, the Schur autoinv works very well in terms of time but has the value function slightly different from others.

In this case the Null space method is the slowest one because of the linear system it has to solve. We must consider that the iterative method (GMRES) doesn't converge to the desired tolerance $1e-12$, the iterate returned has relative residual $3.8e-08$.

The LDL' decomposition is faster than LU because it exploits the symmetry of K .

5.4 Case 4

This case is tested with $n = 10^5$ and $k = 500$. The table shows the value of the function and the *Time* in seconds.

$n = 10^5, k = 100$											
LU		LDL'		GMRES		Schur		Schur autoinv		Null space	
$f(x)$	<i>Time</i>	$f(x)$	<i>Time</i>	$f(x)$	<i>Time</i>	$f(x)$	<i>Time</i>	$f(x)$	<i>Time</i>	$f(x)$	<i>Time</i>
500.00	73.5240	500.00	16.0290	500.00	2.1206	Out of mem- ory	/	502.76	0.6971	500	3.7602

Table 4: Comparison between the methods.

For this case the matrix Q has dimension $10^5 \times 10^5$ and MATLAB can't compute the inverse matrix because it's too large, the Schur autoinv works very well in terms of time but has the value function slightly different from others.

We must consider that the iterative method (GMRES) doesn't converge to the desired tolerance $1e-12$, the iterate returned has relative residual $1.9e-07$.

The LDL' decomposition is faster than LU because it exploits the symmetry of K .

6 Conclusion

The direct methods LU and LDL' are not the best choice because they are affected from the fill-in problem. The figures below show the output of *spy* command for K , L , U matrices for case 1 with the LU decomposition.




Figure 1: Matrix K for the case with $n = 10^4$ and $k = 100$.




Figure 2: Matrix L for the case with $n = 10^4$ and $k = 100$.




Figure 3: Matrix U for the case with $n = 10^4$ and $k = 100$.

We can observe that the number of nonzero elements of K is 49998 out of $(n+m) \times (n+m) = 102010000$, the percentage is less than 0.05%, the number of nonzero elements of L is 1015248 and of U is 1505100, together almost 2.47%.

The GMRES method is not affected from the fill-in problem but doesn't converge within the maximum number of iteration.

The Schur method needs to compute the inverse matrix Q^{-1} and this computation may lead to problem of storage, we see that with $n = 10^5$ this method doesn't work. On the other hand the Schur method `autoinv`, which uses the *backslash* operator insted of computing the inverse matrix, works well with all the cases.

The Null Space method has precise solution but when k is too small w.r.t. n the running time is bigger than others.

In conclusion, to solve a quadratic programming problem with equality constraints only, we prefer the reduced form methods for better results.

Appendix

Schur method

```
1 function [xstar, fxstar, lambda_star] = ...
2     QPeq_Schur(Q, Qinv, c, A, b)
3 % function [xstar, fxstar, lambda_star] = ...
4 %     QPeq_Schur(Q, Qinv, c, A, b)
5 % INPUTS:
6 % Q = Matrix n-by-n of the quadratic loss function
7 % Qinv = inverse of Q
8 % c = n-dimensional vector of the quadratic loss function
9 % A = matrix m-by-n of the equality constraints
10 % b = m-dimensional vector of the equality constraints
11 %
12 % OUTPUTS:
13 % xstar = solution returned by the function
14 % fxstar = value of the loss in xstar
15 % lambda_star = lagrangian multiplier computed by the function
16 %
17
18 % Schur complement computation
19 S = A * Qinv * A';
20
21 % find lambda_star solving the linear system
22 beta = -b - A * Qinv * c;
23 lambda_star = S \ beta;
24
25 % Compute xstar given lambda_star
26 xstar = Qinv * (-c - A' * lambda_star);
27
28 % compute fxstar
29 fxstar = 0.5 * xstar' * Q * xstar + c' * xstar;
30
31 end
```

Schur autoinv method

```
1 function [xstar, fxstar, lambda_star] = ...
2     QPeq_Schur_autoinv(Q, c, A, b)
3 % function [xstar, fxstar, lambda_star] = ...
4 %     QPeq_Schur_autoinv(Q, c, A, b)
5 % INPUTS:
6 % Q = Matrix n-by-n of the quadratic loss function
7 % c = n-dimensional vector of the quadratic loss function
8 % A = matrix m-by-n of the equality constraints
9 % b = m-dimensional vector of the equality constraints
10 %
11 % OUTPUTS:
12 % xstar = solution returned by the function
13 % fxstar = value of the loss in xstar
14 % lambda_star = lagrangian multiplier computed by the function
15 %
16
17 % DON'T USE inv(Q)
18
19 % Schur complement computation
20 S = A * (Q \ A') ;
```

```

21
22 % find lambda_star solving the linear system
23 beta = -b - A * (Q\c);
24 lambda_star = S\beta;
25
26 % Compute xstar given lambda_star
27 xstar = Q\(-c - A' * lambda_star);
28
29 % compute fxstar
30 fxstar = 0.5 * xstar' * Q * xstar + c' * xstar;
31
32
33 end

```

Null Space method

```

1 function [xstar, fxstar, lambda_star, v_star] = ...
2   QPeq_Null(Q, c, A, b, x2)
3 % function [xstar, fxstar, lambda_star, v_star] = ...
4 %   QPeq_Null(Q, c, A, b, x2)
5 % INPUTS:
6 % Q = Matrix n-by-n of the quadratic loss function
7 % c = n-dimensional vector of the quadratic loss function
8 % A = matrix m-by-n of the equality constraints
9 % b = m-dimensional vector of the equality constraints
10 % x2 = (n-m) column vector
11 %
12 % OUTPUTS:
13 % xstar = solution returned by the function
14 % fxstar = value of the loss in xstar
15 % lambda_star = lagrangian multiplier computed by the function
16 % v_star = solution of linear system
17
18
19 % Z initialization
20 [m, n] = size(A);
21 A1 = A(:, 1:m);
22 A2 = A(:, m+1:end);
23 Z = [-A1\A2; speye(n-m)];
24
25 % xhat initialization
26 xhat = [A1\((b - A2 * x2)); x2];
27
28 % compute v_star as solution of:
29 % (Z' Q Z)v = -Z'(Q xhat + c)
30 V = Z' * Q * Z;
31 beta = -Z' * (Q * xhat + c);
32 v_star = V\beta;
33
34 % compute xstar, given v_star and xhat
35 xstar = Z * v_star + xhat;
36
37 % compute lambda_star given xstar
38 lambda_star = (A * A')\(-A * (c + Q * xstar));
39
40 % compute fxstar
41 fxstar = 0.5 * xstar' * Q * xstar + c' * xstar;

```



```

42
43 end

```

LU,LDL',GMRES methods

```

1 %% LU
2 tic
3 %LU decomposition
4 [L1, U] = lu(K);
5 %solve Kw = d
6 wLU = U\ (L1\d);
7 %recall that w = [x;lambda]
8 xLU = wLU(1:n);
9 lambdaLU = wLU(n+1:end);
10 fxLU = 0.5*xLU'*Q*xLU + c'*xLU
11 toc
12
13
14 %% LDL'
15 tic
16 %LDL decomposition
17 [L2,D,P,S] = ld1(K);
18 %solve Kw = d
19 wLDL = S*P*((L2*D*L2') \ (P'*S*d));
20
21 xLDL = wLDL(1:n);
22 lambdaLDL = wLDL(n+1:end);
23 fxLU = 0.5*xLDL'*Q*xLDL + c'*xLDL
24 toc
25
26
27 %% GMRES
28 tic
29 tol = 1e-12;
30 imax = 100; %number of max iterations
31
32 %solve Kw = d
33 wGM = gmres(K, d, [], tol, imax);
34
35 xGM = wGM(1:n);
36 lambdaGM = wGM(n+1:end);
37 fxGM = 0.5*xGM'*Q*xGM + c'*xGM
38 toc

```

Initialization

```

1 n = 10000;
2 k = 100;
3 c = ones (n,1) ;
4 b = ones (k,1);
5 Q = spdiags ([-c 2*c -c] , -1:1 ,n, n );
6 A = repmat (speye(k),1,n/k) ;
7
8 [m,n] = size(A);
9
10
11 O = zeros(m, m);
12 K = [Q A';

```

```

13     A O];
14 d = [-c; b];
15
16 Qin = inv(Q);

```

Testing example

```

1 %Schur with inverse
2
3 disp('**** SCHUR (with Qin) ****')
4 tic
5 [xstar, fxstar, lambda_star] = ...
6     QPeq_Schur(Q, Qin, c, A, b);
7 toc
8 disp('*****')
9 disp(['xstar: ', mat2str(xstar), ';'])
10 disp(['f(xstar): ', num2str(fxstar), ';'])
11 disp(['lambda_star: ', mat2str(lambda_star), ';'])
12 disp('*****')
13
14 %Schur without inverse
15
16 disp('**** SCHUR (without Qin) ****')
17 tic
18 [xstar, fxstar, lambda_star] = ...
19     QPeq_Schur_autoinv(Q, c, A, b);
20 toc
21 disp('*****')
22 disp(['xstar: ', mat2str(xstar), ';'])
23 disp(['f(xstar): ', num2str(fxstar), ';'])
24 disp(['lambda_star: ', mat2str(lambda_star), ';'])
25 disp('*****')
26
27 %Null space
28
29 disp('**** NULL SPACE ****')
30 tic
31 [xstar, fxstar, lambda_star, v_star] = ...
32     QPeq_Null(Q, c, A, b, zeros(n-m, 1));
33 toc
34 disp('*****')
35 disp(['xstar: ', mat2str(xstar), ';'])
36 disp(['f(xstar): ', num2str(fxstar), ';'])
37 disp(['lambda_star: ', mat2str(lambda_star), ';'])
38 disp(['v_star: ', mat2str(v_star), ';'])
39 disp('*****')

```